

digipodium

Machine Learning



Lecture 18

Deep learning

Convolutional Neural Networks

- What are Convolutional Neural Networks?
 - Step 1 - Convolution Operation
 - Step 1(b) - ReLU Layer
 - Step 2 - Pooling
 - Step 3 - Flattening
 - Step 4 - Full Connection
 - Summary
-
- EXTRA: Softmax & Cross-Entropy

Convolutional Neural Networks



WHAT DO YOU SEE HERE?


Examples from the test set (with the network's guesses)



Examples from the test set (with the network's guesses)




Convolutional Neural Networks




Convolutional Neural Networks


B / W Image 2x2px




2d array



Colored Image 2x2px





3d array





Convolutional Neural Networks



Convolutional Neural Networks



Convolutional Neural Networks



0	0	0	0	0	0	0
0	1	0	0	0	1	0
0	0	0	0	0	0	0
0	0	0	1	0	0	0
0	1	0	0	0	1	0
0	0	1	1	1	0	0
0	0	0	0	0	0	0

STEP 1: Convolution



STEP 2: Max Pooling



STEP 3: Flattening



STEP 4: Full Connection

Step 1 - Convolution

0	0	0	0	0	0	0
0	1	0	0	0	1	0
0	0	0	0	0	0	0
0	0	0	1	0	0	0
0	1	0	0	0	1	0
0	0	1	1	1	0	0
0	0	0	0	0	0	0

Input Image

0	0	1
1	0	0
0	1	1

Feature
Detector

Step 1 - Convolution

$$(f * g)(t) \stackrel{\text{def}}{=} \int_{-\infty}^{\infty} f(\tau) g(t - \tau) d\tau$$

0	0	0	0	0	0	0
0	1	0	0	0	1	0
0	0	0	0	0	0	0
0	0	0	1	0	0	0
0	1	0	0	0	1	0
0	0	1	1	1	0	0
0	0	0	0	0	0	0

Input Image



0	0	1
1	0	0
0	1	1


Feature
Detector



0				

Feature Map

Step 1 – Convolution



Step 1 - Convolution

0	0	0	0	0	0	0
0	1	0	0	0	1	0
0	0	0	0	0	0	0
0	0	0	1	0	0	0
0	1	0	0	0	1	0
0	0	1	1	1	0	0
0	0	0	0	0	0	0

Input Image



0	0	1
1	0	0
0	1	1

Feature
Detector



0	1			

Feature Map

Step 1 - Convolution

0	0	0	0	0	0	0
0	1	0	0	0	1	0
0	0	0	0	0	0	0
0	0	0	1	0	0	0
0	1	0	0	0	1	0
0	0	1	1	1	0	0
0	0	0	0	0	0	0

Input Image



0	0	1
1	0	0
0	1	1

Feature
Detector



0	1	0	0	0

Feature Map

Step 1 - Convolution

0	0	0	0	0	0	0
0	1	0	0	0	1	0
0	0	0	0	0	0	0
0	0	0	1	0	0	0
0	1	0	0	0	1	0
0	0	1	1	1	0	0
0	0	0	0	0	0	0

Input Image



0	0	1
1	0	0
0	1	1

Feature
Detector



0	1	0	0	0
0				

Feature Map

Step 1 - Convolution

0	0	0	0	0	0	0
0	1	0	0	0	1	0
0	0	0	0	0	0	0
0	0	0	1	0	0	0
0	1	0	0	0	1	0
0	0	1	1	1	0	0
0	0	0	0	0	0	0

Input Image



0	0	1
1	0	0
0	1	1

Feature
Detector



0	1	0	0	0
0	1	1		

Feature Map

Step 1 - Convolution

0	0	0	0	0	0	0
0	1	0	0	0	1	0
0	0	0	0	0	0	0
0	0	0	1	0	0	0
0	1	0	0	0	1	0
0	0	1	1	1	0	0
0	0	0	0	0	0	0

Input Image



0	0	1
1	0	0
0	1	1


Feature
Detector



0	1	0	0	0
0	1	1	1	0
1	0	1	2	1
1	4	2	1	0
0	0	1	2	1

Feature Map

Step 1 - Convolution



Input Volume (+pad 1) (7x7x3)

 $x[:, :, 0]$

0	0	0	0	0	0	0
0	0	0	1	0	2	0
0	1	0	2	0	1	0

0	1	0	2	2	0	0
0	2	0	0	2	0	0
0	2	1	2	2	0	0
0	0	0	0	0	0	0

 $x[:, :, 1]$

0	0	0	0	0	0	0
0	2	1	2	1	1	0
0	2	1	2	0	1	0

0	0	2	1	0	1	0
0	1	2	2	2	2	0
0	0	1	2	0	1	0
0	0	0	0	0	0	0

 $x[:, :, 2]$

0	0	0	0	0	0	0
0	2	1	1	2	0	0
0	1	0	0	1	0	0

0	0	1	0	0	0	0
0	1	0	2	1	0	0
0	2	2	1	1	1	0
0	0	0	0	0	0	0

Filter W0 (3x3x3)

 $w0[:, :, 0]$

-1	0	1
0	0	1
1	-1	1

 $w0[:, :, 1]$

-1	0	1
1	-1	1
0	1	0

 $w0[:, :, 2]$

-1	1	1
1	1	0
0	-1	0

Bias b0 (1x1x1)

 $b0[:, :, 0]$

1

Filter W1 (3x3x3)

 $w1[:, :, 0]$

0	1	-1
0	-1	0
0	-1	1

 $w1[:, :, 1]$

-1	0	0
1	-1	0
1	-1	0

 $w1[:, :, 2]$

-1	1	-1
0	-1	-1
1	0	0

Bias b1 (1x1x1)

 $b1[:, :, 0]$

0

Output Volume (3x3x2)

 $o[:, :, 0]$

2	3	3
3	7	3
8	10	-3


 $o[:, :, 1]$

-8	-8	-3
-3	1	0
-3	-8	-5

Step 1 – Convolution

Sharpen:

0	0	0	0	0
0	0	-1	0	0
0	-1	5	-1	0
0	0	-1	0	0
0	0	0	0	0




Blur:

0	0	0	0	0
0	1	1	1	0
0	1	1	1	0
0	1	1	1	0
0	0	0	0	0



Edge Enhance:

	0	0	0	
	-1	1	0	
	0	0	0	



Edge Detect:

	0	1	0	
	1	-4	1	
	0	1	0	



Step 1 - Convolution




*

1	0	-1
2	0	-2
1	0	-1




Step 1(B) - ReLU Layer



Step 1 (B) – ReLU Layer



Normal image



Convolution filter



Relu function

Step 2 - Max Pooling



Step 2 - Max Pooling



Image Source: Wikipedia

Step 2 - Max Pooling

0	1	0	0	0
0	1	1	1	0
1	0	1	2	1
1	4	2	1	0
0	0	1	2	1

Feature Map

Max Pooling



1		

Pooled Feature Map

Step 2 - Max Pooling

0	1	0	0	0
0	1	1	1	0
1	0	1	2	1
1	4	2	1	0
0	0	1	2	1

Feature Map

Max Pooling



1	1	

Pooled Feature Map

Step 2 - Max Pooling

0	1	0	0	0
0	1	1	1	0
1	0	1	2	1
1	4	2	1	0
0	0	1	2	1

Feature Map

Max Pooling

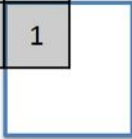


1	1	0

Pooled Feature Map

Step 2 - Max Pooling

0	1	0	0	0
0	1	1	1	0
1	0	1	2	1
1	4	2	1	0
0	0	1	2	1



Feature Map

Max Pooling



1	1	0
4	2	1
0	2	1

Pooled Feature Map

Step 2 – Max Pooling




Image Source: scs.ryerson.ca/~aharley/vis/conv/flat.html

Step 3 - Flattening

1	1	0
4	2	1
0	2	1


Pooled Feature Map

Flattening



1
1
0
4
2
1
0
2
1

Step 3 - Flattening

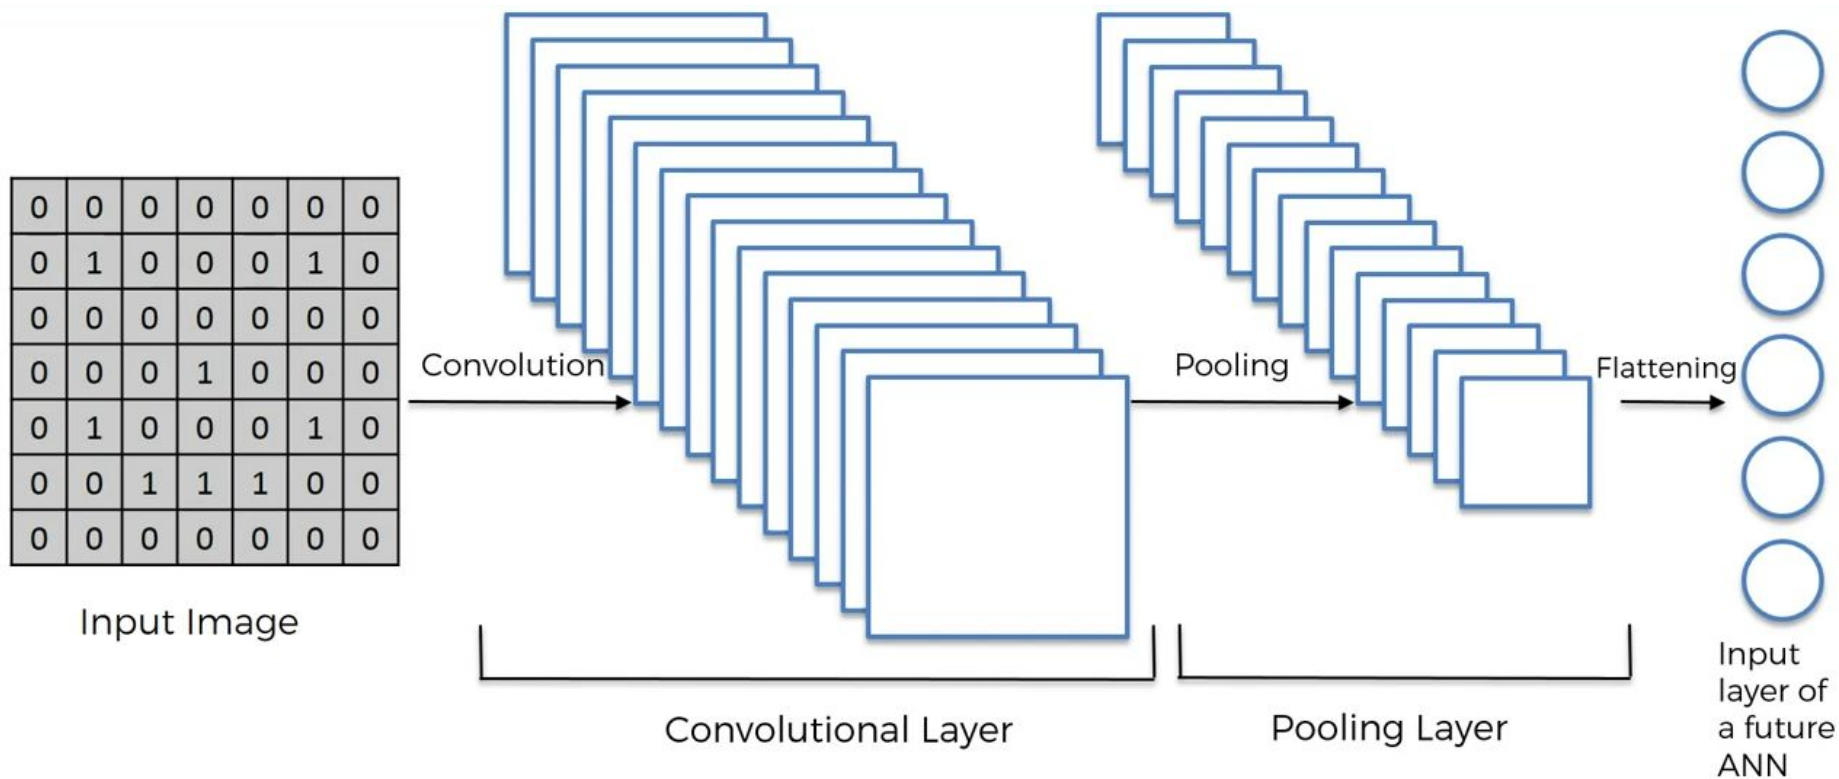


Flattening

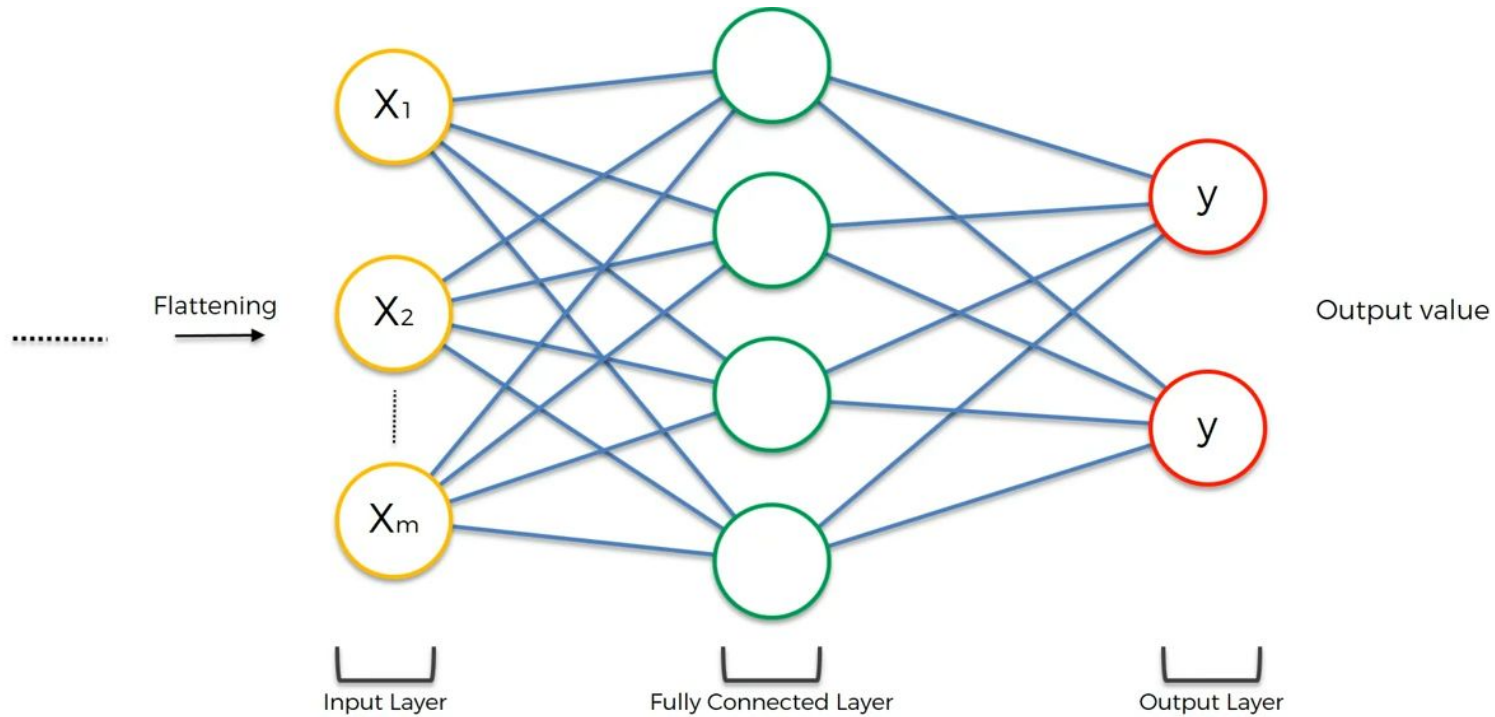


Input layer of a future ANN


Step 3 - Flattening



Step 4 - Full Connection



Step 4 - Full Connection




Step 4 - Full Connection





Flattening

.....




Step 4 – Full Connection








.....
Flattening →




Dog 0.95
Cat 0.05






.....
→ Flattening



Dog → z₁ → 0.95
Cat → z₂ → 0.05

$$f_j(z) = \frac{e^{z_j}}{\sum_k e^{z_k}}$$




Dog 0.9
Cat 0.1

$$H(p, q) = - \sum_x p(x) \log q(x)$$

1
0



Dog

0.9

Cat

0.1

$$H(p, q) = - \sum_x p(x) \log q(x)$$

1

0

NN1

NN2



Dog	1
Cat	0

0.9

0.1

0.6

0.4



Dog	0
Cat	1

0.1

0.9

0.3

0.7



Dog	1
Cat	0

0.4

0.6

0.1

0.9

NN1

NN2

Row	Dog^	Cat^	Dog	Cat
#1	0.9	0.1	1	0
#2	0.1	0.9	0	1
#3	0.4	0.6	1	0

Row	Dog^	Cat^	Dog	Cat
#1	0.6	0.4	1	0
#2	0.3	0.7	0	1
#3	0.1	0.9	1	0

Classification Error

$1/3 = 0.33$

$1/3 = 0.33$

Cross-Entropy

0.381.06

Lower is better



Let's
Code